government_bond_pricing

November 4, 2026

1 Government bond pricing in OSEM

This notebook walks through how OSEM treats a government bond position: from the input file, to projected cash flows, to a calibrated market-consistent price. It is the companion of `bond_pricing.ipynb`, which covers corporate bonds. Both bond types take their cash flows, pricing and calibration from the same class, `FixedCouponBond`, so the mathematics is the same. This notebook concentrates on what is specific to government bonds and on how the two types run together.

On a very high level:

1) Each government bond is saved into a GovBond object.
2) GovBond objects go into the same BondPortfolio as the corporate bonds.
3) Based on the BondPortfolio and market conditions, the cash flows resulting from each bond are reconstructed into a matrix where the cash flow dates are in the columns and individual bonds in rows.
4) Based on the cash flows and the risk-free curve, the z-spread is calculated using the bisection algorithm to return the current market price. For a government bond, that spread is the only place where the sovereign's credit and liquidity risk enter the model.

1.1 Table of contents

1. Methodology
   - Government and corporate bonds compared
   - Notation
   - Step 1 — Cash flow dates
   - Step 2 — Date fractions
   - Step 3 — Cash flow amounts
   - Step 4 — Pricing off the risk-free curve
   - Step 5 — Calibrating the z-spread
   - Step 6 — Calibrating the whole portfolio
   - Step 7 — Use in the full OSEM run
   - What is not modelled
2. Setup
   - Environment and imports
   - Run parameters
3. The government bond portfolio
   - Loading the bonds
   - Which fields drive the price

    - Rows in the wrong file
4. Steps 1 and 3 in numbers: cash flow dates and amounts
    - Coupon and notional flows
    - Cash flow matrices
5. Step 4 in numbers: pricing off the risk-free curve
    - The risk-free curve
    - Pricing state at the modelling date
    - Risk-free price: the gap the z-spread has to close
6. Step 5 in numbers: calibrating one bond
7. Step 6 in numbers: calibrating the whole portfolio
8. Government and corporate bonds in one portfolio
9. Appendix: input data
    - ALM.ini
    - Government_Bond_Portfolio.csv
    - Parameters.csv and the EIOPA files

1.2 1. Methodology

How OSEM turns a government bond description into a price, step by step. The steps are the ones `bond_pricing.ipynb` describes for corporate bonds, because both bond types inherit them from `FixedCouponBond`. No code in this section; the sections that follow run each step on the input portfolio.

1.2.1 Government and corporate bonds compared

	`CorpBond`	`GovBond`
Input file	`Corporate_Bond_Portfolio.csv`	`Government_Bond_Portfolio.csv`
Loader	`get_corporate_bonds`	`get_government_bonds`
`Asset_Type` in the file	`Corporate_Bond`	`Government_Bond`
Issuer field	`nace`, the sector code	`issuer_country`
Credit fields	Sector, country and stress spreads, recovery rate, default probability; loaded but not used	None
Field checks	`_validate_bond_terms`, then recovery rate and default probability	`_validate_bond_terms`, then a non-empty `issuer_country`
Cash flows, pricing, z-spread calibration	Inherited from `FixedCouponBond`	Inherited from `FixedCouponBond`
Portfolio in a run	`BondPortfolio`	The same `BondPortfolio`

`GovBond` has no credit fields. They would add nothing: for a corporate bond they are loaded but play no part in the price, since the calibrated z-spread carries all credit and liquidity risk.

1.2.2 Notation

This table defines the mathematical notation that will be used through the rest of the workbook

Symbol	Meaning
MD	modelling date
VD	valuation date, the date cash flows are discounted to ($VD = MD$ at $t = 0$)
E	end of the modelling window (`settings.end_date`)
I	issue date
M	maturity date
L	last cash flow date, $\min(E, M)$
N	notional amount
c	coupon rate, paid on **each** coupon date
f	coupon frequency (payments/year)
MV	market value (price) at the modelling date
$y(t)$	risk-free spot rate for maturity t, from the EIOPA curve
z	z-spread added to $y(t)$ at every maturity

c is the rate paid on each coupon date, not an annual rate. A bond quoted with an annual coupon c_{ann} paid f times a year is entered with `Coupon_Rate` $= c_{ann}/f$. Italian government bonds (BTPs) pay twice a year, so the 2.45% BTP maturing on 1 September 2050 (IT0005398406) is entered with `Coupon_Rate` 0.01225 and `Frequency` 2.

1.2.3 Step 1 — Cash flow dates

Every bond produces two kinds of cash flow: coupons and the return of notional. A bond still alive when the modelling window closes is redeemed there at par, so both stop at

$$L = \min(E, M)$$

Notional repayment date:

$$t_M^d = L$$

Coupon dates:

$$I < t_1^d < t_2^d < \cdots < t_k^d \leq L$$

The coupon grid is anchored on the issue date and stepped by $12/f$ months. Payments falling before MD are dropped, since only future cash flows are priced. The grid has to land on the maturity date, so in the input file I is the last coupon date on or before the bond's first issue rather than the first issue date itself. A bond first issued between two coupon dates, such as a Bund issued in January that pays every 15 February, would otherwise get its coupons on the wrong day.

Government bonds are often long; 30-year and 50-year issues are common, and Austria has issued 100-year bonds. Four bonds in the input portfolio mature after E: they are repaid at par on E, and their z-spreads are calibrated against those shortened cash flows.

In numbers: section 4.

1.2.4 Step 2 — Date fractions relative to the valuation date

Each remaining cash flow date is converted to a year fraction from the valuation date:

$$t_i = \frac{t_i^d - VD}{365.25}$$

At the modelling date $VD = MD$, which is the case this notebook demonstrates. In a full run the valuation date moves forward with every period and the fractions are re-measured from the new date.

Step 2 runs inside `price_bond` *and has no table of its own.*

1.2.5 Step 3 — Cash flow amounts

Notional cash flow (paid once, at L):

$$cf_M = N$$

Coupon cash flows (paid at each coupon date):

$$cf_i = N \cdot c$$

In numbers: section 4.

1.2.6 Step 4 — Pricing off the risk-free curve

Discounting the cash flows with the risk-free spot curve alone gives

$$MV_{\text{rf}} = \sum_{i=1}^{k} \frac{cf_i}{(1 + y(t_i))^{t_i}} + \frac{cf_M}{(1 + y(t_M))^{t_M}}$$

For a euro area government bond, this curve does not depend on the issuer. EIOPA derives one euro curve from swap rates, less a credit risk adjustment, and publishes it for every euro area member. The `Slovenia` curve this run uses is the euro curve, as section 5 shows, so German, French and Italian bonds are all discounted on the same $y(t)$. Everything that separates them has to come from z in Step 5.

A sovereign can trade richer than the swap curve, as German bonds did in 2023. Its market price is then above MV_{rf}, and the z-spread that closes the gap is negative.

In numbers: section 5.

1.2.7 Step 5 — Calibrating the z-spread to match the market price

OSEM finds a single constant z added to the risk-free yield at every maturity such that the discounted cash flows reproduce the bond's observed market price:

$$MV = \sum_{i=1}^{k} \frac{cf_i}{(1 + y(t_i) + z)^{t_i}} + \frac{cf_M}{(1 + y(t_M) + z)^{t_M}}$$

There is no closed-form solution for z, so OSEM finds it with a **bisection search** (`bisection_spread`) over the bracket $[-0.2, 0.2]$ that `calibrate_bond_portfolio` uses. The bracket includes negative spreads, so a bond priced above MV_{rf} calibrates the same way as one priced below it. If the market price is outside the prices at the two ends of the bracket, `bisection_spread` raises a `ValueError` naming the bond instead of returning a bracket bound. The search stops when the bracket is narrower than `precision` $= 10^{-8}$, a tolerance on the spread rather than on the price.

In numbers: section 6.

1.2.8 Step 6 — Calibrating the whole portfolio

`calibrate_bond_portfolio` repeats Step 5's bisection for every bond in the portfolio, storing one z per asset in `zspread_df`. Each bond is calibrated on its own, so a government bond gets the same z whether it is calibrated alone or together with the corporate bonds.

In numbers: section 7 and section 8.

1.2.9 Step 7 — How this is used in the full OSEM run

This notebook prices bonds once, at $t = 0$. In a full run (`main.py`):

1. `main.py` loads the corporate bond file, then the government bond file if `ALM.ini` names one (`government_bonds` in `[INPUT]`), and keys all the bonds by asset id with `bonds_by_asset_id`. Both types share the bond DataFrames, which are indexed by asset id, so an id used twice, in one file or across the two, stops the run with an error.
2. From then on there is a single `BondPortfolio`, with one set of cash flow matrices (`cpn_df`, `not_df`) and one set of `bd_price_df`, `bd_zspread_df` and `bd_units_df`. The z-spreads are calibrated **once**, at the modelling date. In every later period the bonds are repriced with that *same, fixed* z-spread on the period's projected curve, matured bonds are closed, and `trade()` buys or sells the same share of every holding, government bonds included.
3. `Results.csv` reports coupons, notionals and `Bond market value` for both types together, and `Government bond market value` and `Corporate bond market value` on their own; the two add up to `Bond market value`. The run report draws them as separate bands of the assets chart.

1.2.10 What is not modelled

- **Currency.** A government bond is assumed to be in the currency of the run's curve. A bond in another currency would need its own curve and an exchange-rate model.
- **Default.** There is no default or downgrade model, for government or corporate bonds.
- **Spread stress.** Solvency II exempts bonds of EEA governments in their own currency from spread risk. The run has no spread stress yet, so there is nothing for the exemption to act on.
- **Clean and dirty prices.** `price_bond` discounts the whole next coupon, so it produces a dirty price. Government bonds are quoted clean: enter the clean price plus accrued interest as `Market_Price`, or the z-spread absorbs the accrued interest.
- **Irregular first coupons.** Every coupon is a full period's coupon. A bond issued shortly before MD whose first coupon period is longer or shorter than the rest pays a regular coupon instead.

- **Other structures.** Only fixed-coupon bonds that repay their notional at maturity are supported, so no inflation-linked, floating-rate or callable bonds. The input file leaves them out.
- **New issues.** Matured bonds are not replaced, so the government share of the portfolio falls as the bonds mature.

1.3 2. Setup

Prepares the environment and loads the run parameters. Nothing in this section is specific to bonds.

1.3.1 Environment and imports

Locate the repository root (the folder holding `ALM.ini`), make the `osem` package importable from `src/`, and import the parts of OSEM this notebook uses.

[1]:
```
import os
import sys
import numpy as np
import pandas as pd
from dataclasses import replace
from dateutil.relativedelta import relativedelta
from pathlib import Path

# Walk up to the repository root, the folder holding ALM.ini
base_folder = next(path for path in [Path.cwd(), *Path.cwd().parents] if (path /
 ↪ "ALM.ini").is_file())
sys.path.insert(0, str(base_folder / "src"))
os.chdir(base_folder)
base_folder = os.getcwd()  # Get current working directory

from osem.CurvesClass import Curves
from osem.ImportData import (get_configuration, get_settings, import_SWEiopa,
                             get_corporate_bonds, get_government_bonds)
from osem.BondClasses import BondPortfolio, GovBond, bonds_by_asset_id
from osem.ConfigurationClass import Configuration
from osem.MainLoop import create_cashflow_dataframe, holdings_market_value
```

1.3.2 Run parameters

`ALM.ini` gives the location of every input file, including the government bond file:

[2]:
```
conf: Configuration
conf = get_configuration(os.path.join(base_folder, "ALM.ini"), os)
```

[3]:
```
# Relative to the repository root, which is the working directory
parameters_file = os.path.relpath(conf.input_parameters)
corporate_bond_portfolio_file = os.path.relpath(conf.
 ↪input_corporate_bond_portfolio)
```

```
government_bond_portfolio_file = os.path.relpath(
    conf.input_government_bond_portfolio)
```

The government bond file named in `ALM.ini`:

---

[4]:
```
print(government_bond_portfolio_file)
```

```
Input\Government_Bond_Portfolio.csv
```

The settings object holds data about file locations, information about the run settings and model parameters such as modelling date.

[5]:
```
settings = get_settings(parameters_file)
```

The run that the rest of the notebook prices against:

---

[6]:
```
run_settings = pd.DataFrame(
    # Rendered as strings: in a mixed-type column pandas applies one float␣
 ↪format to every
    # numeric value, which would show a precision of 1e-10 as 0.0
    {"Value": [str(settings.modelling_date),
               str(settings.end_date),
               str(settings.n_proj_years),
               settings.country,
               f"{settings.precision:g}",
               f"{settings.tau:g}"]},
    index=["Modelling date",
           "End of modelling window",
           "Projection years",
           "Country",
           "Precision (curve calibration)",
           "Tau"],
)

display(run_settings)
```

	Value
Modelling date	2023-04-29
End of modelling window	2073-04-29
Projection years	50
Country	Slovenia
Precision (curve calibration)	1e-10
Tau	0.0001

1.4 3. The government bond portfolio

1.4.1 Loading the bonds

`get_government_bonds` reads the government bond file and yields one `GovBond` per row. `bonds_by_asset_id`, the same helper `main.py` uses, keys them by asset id and raises if an id appears twice.

[7]:
```
gov_input = bonds_by_asset_id(
    get_government_bonds(government_bond_portfolio_file))
```

Four bonds serve as examples through the rest of the notebook, each showing one feature that shapes the cash flows:

Asset id	Bond	Feature
DE0001102580	Germany 0% 15/02/2032	Zero coupon: every coupon is 0
FR0010171975	France 4% 25/04/2055	Annual coupon
IT0005398406	Italy 2.45% 01/09/2050	Semi-annual coupon, entered at half the annual rate
SI0002104121	Slovenia 0.6875% 03/03/2081	Matures after the end of the modelling window

The Italian bond as a GovBond:

---

[8]:
```
example_ids = ["DE0001102580", "FR0010171975", "IT0005398406", "SI0002104121"]

display(gov_input["IT0005398406"])
```

```
GovBond(asset_id='IT0005398406', issuer_country='Italy', issue_date=datetime.
↪date(2019, 9, 1), maturity_date=datetime.date(2050, 9, 1), coupon_rate=0.
↪01225, notional_amount=100.0, zspread=0.0, frequency=<Frequency.BIANNUAL: 2>,␣
↪units=25.704, market_price=81.7821)
```

`BondPortfolio` wraps a dictionary of these objects, keyed by asset id, and provides the portfolio-level methods used in the rest of the notebook. It is the same class that holds the corporate bonds.

[9]:
```
gov_portfolio = BondPortfolio(gov_input)
```

The whole government portfolio at a glance. It holds too many bonds to list one by one, so the first table groups them by issuer, and the second shows the four example bonds. `Annual coupon` is derived, not an input: it is `Coupon rate` times `Frequency`. `Market value` is `Units` times `Market price`. Every bond has the same market value, and their total is the first value of the `Government bond market value` column in `Results.csv`.

---

[10]:
```
portfolio_overview = pd.DataFrame(
    [{"Issuer": bond.issuer_country,
      "Maturity": bond.maturity_date,
      "Notional": bond.notional_amount,
      "Coupon rate": bond.coupon_rate,
      "Frequency": int(bond.frequency),
      "Annual coupon": bond.coupon_rate * int(bond.frequency),
      "Units": bond.units,
      "Market price": bond.market_price,
      "Market value": bond.units * bond.market_price}
     for bond in gov_portfolio.bonds.values()],
    index=list(gov_portfolio.bonds),
)

by_issuer = portfolio_overview.groupby("Issuer").agg(**{
    "Bonds": ("Maturity", "size"),
    "First maturity": ("Maturity", "min"),
    "Last maturity": ("Maturity", "max"),
    "Frequency": ("Frequency", "max"),
    "Lowest annual coupon": ("Annual coupon", "min"),
    "Highest annual coupon": ("Annual coupon", "max"),
    "Market value": ("Market value", "sum"),
}).sort_values("Bonds", ascending=False)

display(by_issuer)
print(f"{len(portfolio_overview)} bonds of {len(by_issuer)} governments, "
      f"total market value: {portfolio_overview['Market value'].sum():,.2f}")
display(portfolio_overview.loc[example_ids])
```

Issuer	Bonds	First maturity	Last maturity	Frequency	\
Italy	61	2026-12-01	2072-03-01	2	
Germany	44	2027-02-15	2053-08-15	1	
Spain	39	2027-02-01	2071-11-02	1	
France	39	2026-11-25	2072-05-25	1	
Belgium	27	2027-06-22	2071-06-22	1	
Ireland	25	2027-05-15	2050-05-15	1	
Austria	24	2027-04-20	2120-06-30	1	
Netherlands	19	2027-01-15	2054-01-15	1	
Slovakia	19	2027-01-21	2068-06-12	1	
Finland	17	2027-04-15	2052-04-15	1	
Portugal	15	2027-04-14	2052-04-12	1	
Slovenia	15	2027-03-22	2081-03-03	1	
Greece	14	2027-04-22	2052-01-24	1	
Luxembourg	10	2026-11-13	2043-03-02	1	
Croatia	5	2029-06-19	2041-03-04	1	
Estonia	2	2030-06-10	2032-10-12	1	
Bulgaria	1	2033-01-27	2033-01-27	1	

```
             Lowest annual coupon  Highest annual coupon   Market value
Issuer
Italy                     0.00250                0.06500  128229.911753
Germany                   0.00000                0.06500   92493.721982
Spain                     0.00000                0.06000   81983.107037
France                    0.00000                0.05750   81983.084142
Belgium                   0.00000                0.05500   56757.537797
Ireland                   0.00000                0.05920   52553.259243
Austria                   0.00000                0.06250   50451.140129
Netherlands               0.00000                0.05500   39940.469107
Slovakia                  0.00125                0.04400   39940.471505
Finland                   0.00000                0.02750   35736.215272
Portugal                  0.00300                0.04125   31531.944315
Slovenia                  0.00000                0.03625   31531.951538
Greece                    0.00750                0.04250   29429.826934
Luxembourg                0.00000                0.03250   21021.302939
Croatia                   0.01125                0.02875   10510.646762
Estonia                   0.00125                0.04000    4204.255552
Bulgaria                  0.04500                0.04500    2102.130445

376 bonds of 17 governments, total market value: 790,400.98

                Issuer    Maturity  Notional  Coupon rate  Frequency  \
DE0001102580   Germany  2032-02-15     100.0     0.000000          1
FR0010171975    France  2055-04-25     100.0     0.040000          1
IT0005398406     Italy  2050-09-01     100.0     0.012250          2
SI0002104121  Slovenia  2081-03-03     100.0     0.006875          1

              Annual coupon    Units  Market price  Market value
DE0001102580       0.000000  28.5704       73.5773   2102.132892
FR0010171975       0.040000  19.3020      108.9074   2102.130635
IT0005398406       0.024500  25.7040       81.7821   2102.127098
SI0002104121       0.006875  60.7675       34.5930   2102.130128
```

1.4.2 Which fields drive the price

Seven of the ten `GovBond` fields enter the price. The other three identify the position and size it.

`GovBond` field	CSV column	Used in pricing	Role
`asset_id`	`Asset_ID`	–	Key of the portfolio dictionary and row index of every DataFrame
`issuer_country`	`Issuer_Country`	No	Issuing government; must not be empty. It does not choose the curve: every bond is priced on the curve of `country` in `Parameters.csv`

GovBond field	CSV column	Used in pricing	Role
issue_date	Issue_Date	Yes	Anchors the coupon grid (I, Step 1)
maturity_date	Maturity_Date	Yes	Notional repayment date (M, Step 1)
frequency	Frequency	Yes	Spacing of the coupon dates, $12/f$ months (f, Step 1)
notional_amount	Notional_Amount	Yes	Notional repaid at L (N, Step 3)
coupon_rate	Coupon_Rate	Yes	Coupon paid on each coupon date, as a fraction of N (c, Step 3)
market_price	Market_Price	Yes	Calibration target (MV, Step 5)
zspread	Z_Spread	Starting value only	Seeds zspread_df; overwritten by the calibrated z-spread (Step 6)
units	Units	No	Number of bonds held; scales the cash flows in the main loop, not the per-unit price

1.4.3 Rows in the wrong file

Both bond loaders check the Asset_Type of every row. Pointed at the corporate bond file, get_government_bonds stops at the first row, instead of reading corporate columns as if they were government ones:

---

[11]:
```
try:
    list(get_government_bonds(corporate_bond_portfolio_file))
except ValueError as error:
    print(error)
```

```
Asset BE6285455497 in Input\Corporate_Bond_Portfolio.csv has Asset_Type
Corporate_Bond, but this file can only hold Government_Bond rows
```

1.5 4. Steps 1 and 3 in numbers: cash flow dates and amounts

This section builds the coupon and notional schedules of Step 1 and the amounts of Step 3 for every government bond. Step 2, the conversion to year fractions, happens inside price_bond when the flows are discounted in the next section.

1.5.1 Coupon and notional flows

The coupon payments and the return of the notional are simulated separately.

create_coupon_flows returns the coupon dates and amounts of every bond, as a dictionary keyed by asset id:

[12]:
```
coupon_flows = gov_portfolio.create_coupon_flows(
    modelling_date=settings.modelling_date, end_date=settings.end_date)
```

`create_maturity_flows` returns the notional repayments in the same shape:

[13]:
```
notional_flows = gov_portfolio.create_maturity_flows(
    terminal_date=settings.end_date)
```

A compact view of the same two dictionaries for the four example bonds. The German bond's coupons are all 0. The Italian bond pays twice a year, so it has more coupons than the French bond although it matures five years earlier, each at half its annual rate. The Slovenian bond outlives the modelling window: its notional is repaid on E, and its coupons stop at the last coupon date before it.

---

[14]:
```
coupon_schedule = pd.DataFrame(
    [{"Issuer": gov_input[aid].issuer_country,
      "Coupons": len(coupons),
      "Coupon amount": next(iter(coupons.values())),
      "First coupon": min(coupons),
      "Last coupon": max(coupons),
      "Notional repaid": min(notional_flows[aid])}
     for aid, coupons in coupon_flows.items()],
    index=list(coupon_flows),
)

display(coupon_schedule.loc[example_ids])
```

```
               Issuer  Coupons  Coupon amount First coupon Last coupon  \
DE0001102580   Germany        9         0.0000   2024-02-15  2032-02-15
FR0010171975    France       32         4.0000   2024-04-25  2055-04-25
IT0005398406     Italy       55         1.2250   2023-09-01  2050-09-01
SI0002104121  Slovenia       50         0.6875   2024-03-03  2073-03-03

              Notional repaid
DE0001102580       2032-02-15
FR0010171975       2055-04-25
IT0005398406       2050-09-01
SI0002104121       2073-04-29
```

1.5.2 Cash flow matrices

All cash flows can be represented in a matrix with all possible cash flow dates as columns and all bonds as rows. The non-zero entries are the cash flow per unit of the bond at that date. The columns are the unique payment dates across the portfolio, from `unique_dates_profile`, computed separately for coupons and notional repayments:

[15]:
```
unique_list = gov_portfolio.unique_dates_profile(coupon_flows)
unique_terminal_list = gov_portfolio.unique_dates_profile(notional_flows)
```

A dataframe of cash flows and notional amount payments is created:

[16]:
```
# Dataframe with government bond coupon cash flows
cash_flows = create_cashflow_dataframe(coupon_flows, unique_list)
# Dataframe with government bond notional cash flows
notional_cash_flows = create_cashflow_dataframe(notional_flows,
                                                unique_terminal_list)
```

The full coupon matrix has one row per bond and one column per payment date. Its size, then the coupons the four example bonds pay in the first year of the run; dates on which none of the four pays are left out:

---

[17]:
```
print(f"Coupon matrix: {cash_flows.shape[0]} bonds by {cash_flows.shape[1]}␣
 ↪payment dates")

first_year = [d for d in cash_flows.columns
              if d < settings.modelling_date + relativedelta(years=1)]
example_coupons = cash_flows.loc[example_ids, first_year]
display(example_coupons.loc[:, (example_coupons != 0).any()])
```

```
Coupon matrix: 376 bonds by 2042 payment dates
```

	2023-09-01	2024-03-01	2024-03-03	2024-04-25
DE0001102580	0.000	0.000	0.0000	0.0
FR0010171975	0.000	0.000	0.0000	4.0
IT0005398406	1.225	1.225	0.0000	0.0
SI0002104121	0.000	0.000	0.6875	0.0

The notional matrix in the same way: its size, then the repayment dates of the four example bonds. The Slovenian bond is repaid on E, not on its maturity date:

---

[18]:
```
print(f"Notional matrix: {notional_cash_flows.shape[0]} bonds by "
      f"{notional_cash_flows.shape[1]} repayment dates")

example_notionals = notional_cash_flows.loc[example_ids]
display(example_notionals.loc[:, (example_notionals != 0).any()])
```

```
Notional matrix: 376 bonds by 345 repayment dates
```

	2032-02-15	2050-09-01	2055-04-25	2073-04-29
DE0001102580	100.0	0.0	0.0	0.0
FR0010171975	0.0	0.0	100.0	0.0
IT0005398406	0.0	100.0	0.0	0.0
SI0002104121	0.0	0.0	0.0	100.0

1.6 5. Step 4 in numbers: pricing off the risk-free curve

This section prices the cash flows from section 4 off the risk-free curve (Step 4).

1.6.1 The risk-free curve

`import_SWEiopa()` reads the yield curve of the run's `country` from the EIOPA files, and the `Curves` object calibrates and projects it, exactly as in `main.py`.

[19]:
```
[maturities_country, curve_country, extra_param, Qb] = import_SWEiopa(
    settings.EIOPA_param_file, settings.EIOPA_curves_file, settings.country)
# Curves object with information about term structure
curves = Curves(extra_param["UFR"] / 100, settings.precision, settings.tau,
                settings.modelling_date, settings.country)
```

[20]:
```
curves.set_observed_term_structure(
    maturity_vec=np.asarray(curve_country.index, dtype=np.float64),
    yield_vec=np.asarray(curve_country, dtype=np.float64))
curves.calc_fwd_rates()
# The +1 matches main.py: the last projection period still has to be priced, so␣
 ↪the
# projected curve has to reach one year past the end of the modelling window.
curves.project_forward_rate(settings.n_proj_years+1)
curves.calibrate_projected(settings.n_proj_years+1, 0.05, 0.5, 1000)
```

The run's curve is the euro curve. Reading the `Euro` column of the same file and comparing it with the run's curve at a few maturities, and then at every maturity:

---

[21]:
```
_, euro_curve, _, _ = import_SWEiopa(
    settings.EIOPA_param_file, settings.EIOPA_curves_file, "Euro")

curve_comparison = pd.DataFrame({settings.country: curve_country,
                                 "Euro": euro_curve})
display(curve_comparison.iloc[[0, 4, 9, 19, 29]])
print(f"The {settings.country} curve equals the euro curve at every maturity: "
      f"{bool((curve_country == euro_curve).all())}")
```

```
         Slovenia     Euro
Country
1         0.03271  0.03271
5         0.03488  0.03488
10        0.03577  0.03577
20        0.03597  0.03597
30        0.03514  0.03514
```

```
The Slovenia curve equals the euro curve at every maturity: True
```

1.6.2 Pricing state at the modelling date

`init_bond_portfolio_to_dataframe` sets up the three DataFrames the pricing works on, each with one column for the modelling date: the market price of each bond, its z-spread and the units held.

[22]:
```
[market_price_df, zspread_df, units_df] = (
    gov_portfolio.init_bond_portfolio_to_dataframe(settings.modelling_date))
```

`zspread_df` starts from the `Z_Spread` column of the government bond file, which is 0 for every bond. The calibration in sections 6 and 7 overwrites it. The starting z-spreads of the four example bonds, then a count over the whole portfolio:

---

[23]:
```
display(zspread_df.loc[example_ids])
print(f"Bonds starting from a non-zero z-spread: {int((zspread_df.iloc[:, 0] !=␣
  ↪0).sum())}")
```

	2023-04-29
DE0001102580	0.0
FR0010171975	0.0
IT0005398406	0.0
SI0002104121	0.0

```
Bonds starting from a non-zero z-spread: 0
```

All pricing in this notebook happens at the modelling date, so the projection year is 0:

[24]:
```
proj_period = 0
```

1.6.3 Risk-free price: the gap the z-spread has to close

Pricing the cash flows at $z = 0$ gives MV_{rf} of Step 4. A bond priced below it is cheaper than the euro curve implies and needs a positive z-spread; a bond priced above it needs a negative one. The input file sets every market price to MV_{rf} itself, rounded to four decimals, so there is no gap to close and every z-spread will come out at 0. Observed market prices would instead give each bond its spread over the euro curve. The four example bonds, then the largest gap in the portfolio:

---

[25]:
```
risk_free_price = {}
for asset_id, bond in gov_portfolio.bonds.items():
    risk_free_price[asset_id] = bond.price_bond(
        cash_flows.loc[asset_id], notional_cash_flows.loc[asset_id],
        settings.modelling_date, proj_period, curves, 0.0)

issuer = pd.Series({aid: b.issuer_country for aid, b in gov_input.items()})
market_price = pd.Series({aid: b.market_price for aid, b in gov_input.items()})
price_comparison = pd.DataFrame({
```

```
    "Issuer": issuer,
    "Risk free price": pd.Series(risk_free_price),
    "Market price": market_price,
})
price_comparison["Gap to close"] = (price_comparison["Risk free price"]
                                    - price_comparison["Market price"])

display(price_comparison.loc[example_ids])
print(f"Largest gap in the portfolio: {price_comparison['Gap to close'].abs().
  ↪max():.6f}")
```

	Issuer	Risk free price	Market price	Gap to close
DE0001102580	Germany	73.577270	73.5773	-0.000030
FR0010171975	France	108.907407	108.9074	0.000007
IT0005398406	Italy	81.782123	81.7821	0.000023
SI0002104121	Slovenia	34.593003	34.5930	0.000003

```
Largest gap in the portfolio: 0.000050
```

1.7 6. Step 5 in numbers: calibrating one bond

Step 5 for the Italian bond, IT0005398406. At its file price the bisection would simply return 0. To show it finding a spread, this section calibrates the bond as if it traded 5 below its file price, as it would if the market asked for a premium over the euro curve. The lower price needs a positive spread, and the bracket $[-0.2, 0.2]$ includes it. The bond at the lower price is a copy made with `replace`, so the portfolio is unchanged.

[26]:
```
italian_bond = gov_portfolio.bonds["IT0005398406"]
cheaper_bond = replace(italian_bond, market_price=italian_bond.market_price - 5)

calibrated_spread = cheaper_bond.bisection_spread(x_start=-0.2
                                , x_end=0.2
                                , modelling_date=settings.modelling_date
                                , end_date=settings.end_date
                                , proj_period=proj_period
                                , curves=curves
                                , precision=0.00000001
                                , max_iter=100000)
```

The calibrated z-spread, as a decimal and in basis points:

---

[27]:
```
print(f"{calibrated_spread:.6f} ({calibrated_spread * 10000:.1f} bp)")
```

```
0.003457 (34.6 bp)
```

Discounting the bond's cash flows at the calibrated z-spread gives back the lower market price:

[28]:
```
calibrated_bond = cheaper_bond.price_bond(
    cash_flows.loc["IT0005398406"], notional_cash_flows.loc["IT0005398406"],
    settings.modelling_date, proj_period, curves, calibrated_spread)
```

---

[29]:
```
print(f"Repriced value {calibrated_bond:.6f}, "
      f"market price {cheaper_bond.market_price:.4f}")
```

```
Repriced value 76.782106, market price 76.7821
```

1.8 7. Step 6 in numbers: calibrating the whole portfolio

`calibrate_bond_portfolio` repeats the bisection for every bond (Step 6) and writes the results into `zspread_df`:

[30]:
```
zspread_df = gov_portfolio.calibrate_bond_portfolio(zspread_df, settings,
                                                    proj_period, curves)
```

---

[31]:
```
display(zspread_df.loc[example_ids])
print(f"Largest calibrated z-spread, in absolute value: {zspread_df.iloc[:, 0].
  ↪abs().max():.2e}")
```

```
                2023-04-29
DE0001102580 -5.364418e-08
FR0010171975  5.960464e-09
IT0005398406  1.788139e-08
SI0002104121  5.960464e-09
```

```
Largest calibrated z-spread, in absolute value: 1.49e-07
```

Repricing the portfolio with the calibrated spreads reproduces the market prices. The residual is the price cost of stopping the bisection once the spread bracket is narrower than `precision`, as described in Step 5.

Every calibrated z-spread is 0, to within that precision, because the market prices in the input file are the risk-free prices of section 5. With observed prices instead, the spreads would follow the euro government bond market: German bonds below the swap-based euro curve, Italian bonds furthest above it. The four example bonds, then the largest residual in the portfolio:

---

[32]:
```
calibrated_price_df = gov_portfolio.price_bond_portfolio(
    cash_flows, notional_cash_flows, settings, proj_period, curves,
    zspread_df, market_price_df.copy(), settings.modelling_date)

calibration_check = pd.DataFrame({
```

```
    "Issuer": issuer,
    "Calibrated z-spread": zspread_df.iloc[:, 0],
    "Z-spread (bp)": (zspread_df.iloc[:, 0] * 10000).round(1),
    "Repriced value": calibrated_price_df.iloc[:, 0],
    "Market price": market_price,
})
calibration_check["Residual"] = (calibration_check["Repriced value"]
                                 - calibration_check["Market price"])

display(calibration_check.loc[example_ids])
print(f"Largest residual in the portfolio: {calibration_check['Residual'].abs().
  ↪max():.2e}")
```

```
                Issuer  Calibrated z-spread  Z-spread (bp)  Repriced value  \
DE0001102580   Germany        -5.364418e-08           -0.0       73.577303
FR0010171975    France         5.960464e-09            0.0      108.907395
IT0005398406     Italy         1.788139e-08            0.0       81.782096
SI0002104121  Slovenia         5.960464e-09            0.0       34.592996

              Market price  Residual
DE0001102580       73.5773  0.000003
FR0010171975      108.9074 -0.000005
IT0005398406       81.7821 -0.000004
SI0002104121       34.5930 -0.000004

Largest residual in the portfolio: 1.18e-05
```

1.9 8. Government and corporate bonds in one portfolio

`main.py` keeps no separate government portfolio. It loads both files and keys every bond by asset id in one dictionary, as Step 7 describes, so each step above runs once for both types:

[33]:
```
bonds = list(get_corporate_bonds(corporate_bond_portfolio_file))
bonds += get_government_bonds(government_bond_portfolio_file)
bond_portfolio = BondPortfolio(bonds_by_asset_id(bonds))
gov_ids = bond_portfolio.asset_ids_of(GovBond)

[all_price_df, all_zspread_df, all_units_df] = (
    bond_portfolio.init_bond_portfolio_to_dataframe(settings.modelling_date))
all_zspread_df = bond_portfolio.calibrate_bond_portfolio(
    all_zspread_df, settings, proj_period, curves)
```

The combined portfolio by bond type. The government bonds are about 70% of the bond holdings by market value. Their z-spreads are all 0, while the corporate ones range from about -1% to +1%:

---

[34]:
```
bond_type = pd.Series({aid: type(bond).__name__
                       for aid, bond in bond_portfolio.bonds.items()})
market_value = all_units_df.iloc[:, 0] * all_price_df.iloc[:, 0]
all_spreads = all_zspread_df.iloc[:, 0]

by_type = pd.DataFrame({
    "Bonds": bond_type.value_counts(),
    "Market value": market_value.groupby(bond_type).sum(),
    "Lowest z-spread": all_spreads.groupby(bond_type).min(),
    "Highest z-spread": all_spreads.groupby(bond_type).max(),
})

display(by_type)
```

	Bonds	Market value	Lowest z-spread	Highest z-spread
CorpBond	201	330399.990020	-1.046994e-02	1.019648e-02
GovBond	376	790400.976452	-1.490116e-07	1.490116e-07

Each bond is calibrated on its own, so the government bonds keep the spreads of section 7. Their market value, taken from the shared DataFrames at the ids `asset_ids_of(GovBond)` lists, as `main.py` does, is the opening value of `Government bond market value` in `Results.csv`:

---

[35]:
```
same_spreads = bool((all_zspread_df.loc[gov_ids] == zspread_df).all().all())
government_value = holdings_market_value(all_price_df.loc[gov_ids],
                                         all_units_df.loc[gov_ids],
                                         settings.modelling_date)

print(f"Same z-spreads as when calibrated alone: {same_spreads}")
print(f"Government bond market value at the modelling date: "
      f"{government_value:,.2f}")
```

```
Same z-spreads as when calibrated alone: True
Government bond market value at the modelling date: 790,400.98
```

Because the two types share these DataFrames, an asset id may be used only once across both files. A government bond given the id of a corporate bond stops the run:

---

[36]:
```
clash = replace(gov_input["SI0002104121"], asset_id=bonds[0].asset_id)
try:
    bonds_by_asset_id([*bonds, clash])
except ValueError as error:
    print(error)
```

```
Asset ID BE6285455497 is used by two bonds (CorpBond and GovBond). Bond asset
IDs must be unique across the corporate and government bond input files.
```

1.10 9. Appendix: input data

The input files this notebook reads, all in the `Input` folder. Their locations come from the `[INPUT]` section of `ALM.ini`, except the EIOPA files, whose paths are set in `Parameters.csv`. Dates are written as `dd/mm/yyyy`.

1.10.1 ALM.ini

Two keys in the `[INPUT]` section name the bond files:

Key	Example	Meaning
`corporate_bonds`	`Corporate_Bond_Portfolio.csv`	Corporate bond file, in the input folder. An `ALM.ini` written before government bonds were added calls this key `bonds`
`government_bonds`	`Government_Bond_Portfolio.csv`	Government bond file, in the input folder. Optional: without it a run holds no government bonds, `Government bond market value` is 0, and an `ALM.ini` written before government bonds were added still loads

1.10.2 Government_Bond_Portfolio.csv

The government bond portfolio, one row per bond position. Loaded by `get_government_bonds`, one `GovBond` per row. Section 3 shows which of these fields enter the price. The positions are real fixed-coupon bonds of 17 euro area governments, taken from the ECB list of eligible marketable assets of 3 November 2026: every nominal bond of those governments that was issued by the modelling date and is still outstanding. `Market_Price` is the price at a zero z-spread on the run's curve, so every bond calibrates to a z-spread of 0. The examples are the Italian bond's row.

Column	Example	Meaning
`Asset_ID`	`IT0005398406`	Unique id of the position, read as text: the bond's ISIN. No corporate bond may use it either
`Asset_Type`	`Government_Bond`	Asset type; `get_government_bonds` raises if it is not `Government_Bond`
`Issuer_Country`	`Italy`	Issuing government; must not be empty
`Issue_Date`	`01/09/2019`	Issue date I; anchors the coupon grid. In this file it is the last coupon date on or before the bond's first issue, so the grid lands on the maturity date

Column	Example	Meaning
`Maturity_Date`	`01/09/2050`	Maturity date M; the notional is repaid on it
`Notional_Amount`	`100`	Notional amount N
`Coupon_Rate`	`0.01225`	Coupon c paid on **each** coupon date, as a fraction of the notional. Not annualised: this bond pays 2.45% a year in two coupons
`Z_Spread`	`0`	Starting z-spread; replaced by the calibrated value
`Frequency`	`2`	Coupon payments per year f: 1, 2, 3, 4 or 12
`Units`	`25.7040`	Number of bonds held
`Market_Price`	`81.7821`	Price of one bond at the modelling date, on the same scale as `Notional_Amount`, including accrued interest; the calibration target MV

1.10.3 Parameters.csv and the EIOPA files

`Parameters.csv`, `Curves_no_VA.csv` and `Param_no_VA.csv` are read exactly as for corporate bonds; the appendix of `bond_pricing.ipynb` describes their columns. Of the parameters, this notebook uses `Modelling_Date`, `n_proj_years`, `country`, `Precision`, `Tau` and the paths of the two EIOPA files.